

IMAGE CLASSIFICATION

Github link to the jupyter notebook:

https://github.com/ashutosh-999/img_cls_in.git

SCREENSHOTS

1. Neural Network model:

```
class TinyVGG(nn.Module):
    def __init__(self, input_shape, hidden_units, output_shape):
        super().__init__()
        self.conv_block_1 = nn.Sequential(
            nn.Conv2d(in_channels=input_shape,
                      out_channels=hidden_units,
                      kernel_size=3,
                      stride=1,
                      padding=1),
            nn.ReLU(),
            nn.Conv2d(in_channels=hidden_units,
                      out_channels=hidden_units,
                      kernel_size=3,
                      stride=1,
                      padding=1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2)
        )
        self.conv_block_2 = nn.Sequential(
            nn.Conv2d(in_channels=hidden_units,
                      out_channels=hidden_units,
                      kernel_size=3,
                      stride=1,
                      padding=1),
            nn.ReLU(),
            nn.Conv2d(in_channels=hidden_units,
                      out_channels=hidden_units,
                      kernel_size=3,
                      stride=1,
                      padding=1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2)
        )
        self.classifier = nn.Sequential(
            nn.Flatten(),
            nn.Linear(in_features=hidden_units*16*16,
                      out_features=output_shape))

    def forward(self, x):
        x = self.conv_block_1(x)
        # print(f"Layer 1 shape: {x.shape}")
        x = self.conv_block_2(x)
        # print(f"Layer 2 shape: {x.shape}")
        x = self.classifier(x)
        # print(f"Layer 3 shape: {x.shape}")
        return x
```

2. Downloading the dataset:

```
# Setup path to data folder
data_path = Path("data/")
image_path = data_path / "pizza_steak_sushi_20_percent"

# If the image folder doesn't exist, download it and prepare it...
if image_path.is_dir():
    print(f"{image_path} directory exists.")
else:
    print(f"Did not find {image_path} directory, creating one...")
    image_path.mkdir(parents=True, exist_ok=True)

# Download pizza, steak, sushi data
with open(data_path / "pizza_steak_sushi_20_percent.zip", "wb") as f:
    request = requests.get("https://github.com/mrdbourke/pytorch-deep-learning/raw/main/data/pizza_steak_sushi_20_percent.zip")
    print("Downloading pizza, steak, sushi 20% data...")
    f.write(request.content)

# Unzip pizza, steak, sushi data
with zipfile.ZipFile(data_path / "pizza_steak_sushi_20_percent.zip", "r") as zip_ref:
    print("Unzipping pizza, steak, sushi 20% data...")
    zip_ref.extractall(image_path)
```

3. Optimizer and Loss Function used:

```
loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model_4.parameters(), lr=0.001)
```

4. Accuracy of the model:

Epoch: 34		train_loss: 0.0039		train_acc: 1.0000		test_loss: 2.8498		test_acc: 0.5159
Epoch: 35		train_loss: 0.0020		train_acc: 1.0000		test_loss: 2.9435		test_acc: 0.5068
Epoch: 36		train_loss: 0.0016		train_acc: 1.0000		test_loss: 2.9854		test_acc: 0.5068
Epoch: 37		train_loss: 0.0013		train_acc: 1.0000		test_loss: 3.0201		test_acc: 0.5159
Epoch: 38		train_loss: 0.0012		train_acc: 1.0000		test_loss: 3.0623		test_acc: 0.5159
Epoch: 39		train_loss: 0.0011		train_acc: 1.0000		test_loss: 3.1062		test_acc: 0.5222
Epoch: 40		train_loss: 0.0010		train_acc: 1.0000		test_loss: 3.1454		test_acc: 0.5222
Epoch: 41		train_loss: 0.0009		train_acc: 1.0000		test_loss: 3.1671		test_acc: 0.5159
Epoch: 42		train_loss: 0.0008		train_acc: 1.0000		test_loss: 3.1913		test_acc: 0.5284
Epoch: 43		train_loss: 0.0007		train_acc: 1.0000		test_loss: 3.2148		test_acc: 0.5222
Epoch: 44		train_loss: 0.0007		train_acc: 1.0000		test_loss: 3.2454		test_acc: 0.5284
Epoch: 45		train_loss: 0.0007		train_acc: 1.0000		test_loss: 3.2679		test_acc: 0.5375
Epoch: 46		train_loss: 0.0006		train_acc: 1.0000		test_loss: 3.2939		test_acc: 0.5312
Epoch: 47		train_loss: 0.0006		train_acc: 1.0000		test_loss: 3.3156		test_acc: 0.5222
Epoch: 48		train_loss: 0.0006		train_acc: 1.0000		test_loss: 3.3332		test_acc: 0.5222
Epoch: 49		train_loss: 0.0005		train_acc: 1.0000		test_loss: 3.3536		test_acc: 0.5159
Epoch: 50		train_loss: 0.0005		train_acc: 1.0000		test_loss: 3.3776		test_acc: 0.5159